

Generative AI Assignment 3 Report

<u>Activation Function</u>	<u>Accuracy</u>	<u>Exec Time (seconds)</u>
Relu	51.59%	28.48s
Tanh	45.70%	28.16s
Identity	39.52%	27.96s

<u>Optimizers</u>	<u>Accuracy</u>	<u>Exec Time</u>
SGD	52.14%	37.92s
ADAM	52.37%	47.70s
RMSPROP	51.58%	43.50s

<u>Regularization</u>	<u>Accuracy</u>	<u>Exec Time</u>
Baseline	49.46%	156.44s
Dropout	40.25%	157.14s
Batch Norm	51.09%	158.06s
Weight Int	47.28%	156.45s

Generative AI Assignment 3 Report

Activation Function

1. **ReLU**: Achieved the highest accuracy of 51.59% and reasonable execution time, indicating it is the most effective activation function for this model.
2. **Tanh**: Recorded a lower accuracy of 45.70%, suggesting it is less effective than ReLU in this context, despite a similar execution time.
3. **Identity (Linear)**: Obtained the lowest accuracy of 39.52% with the fastest execution time, highlighting its inadequacy for this task compared to the other activation functions.

Optimizers:

1. **SGD**: Achieved 52.14% accuracy with the fastest execution time, indicating a solid choice for quick training but with room for improvement in accuracy.
2. **Adam**: Attained the highest accuracy of 52.37%, demonstrating its effectiveness in adapting learning rates for better convergence.
3. **RMSprop**: Delivered 51.58% accuracy, performing well but slightly behind SGD and Adam, suggesting it may not be as well-suited for this task.

Regularization:

1. **BatchNorm**: seems to provide the best improvement in accuracy without significantly impacting execution time. This suggests that stabilizing the internal activations had the most beneficial effect.

Generative AI Assignment 3 Report

2. **Dropout:** seems to hurt the model's performance in this case, likely because the model's complexity wasn't high enough to justify heavy regularization.
3. **Weight Initialization:** didn't have a large impact here, but it might still be useful in larger or deeper networks.